

Problem1 Implement and improve BaseNet on CIFAR-10

1. Implement BaseNet:

BaseNet

1.1 Code:

```
class BaseNet(nn.Module):
    def __init__(self):
        super(BaseNet, self).__init__()

        # TODO: define your model here
        # pass
        # torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0,
dilation=1, groups=1, bias=True, padding_mode='zeros')
        # torch.nn.MaxPool2d(kernel_size, stride=None, padding=0, dilation=1,
return_indices=False, ceil_mode=False)
        # torch.nn.Linear(in_features, out_features, bias=True)

        self.conv1 = nn.Conv2d(3, 6, 5)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.conv3 = nn.Conv2d(16,64,3)
        self.pool = nn.MaxPool2d(2, 2)

        self.fc_net = nn.Sequential(
            nn.Linear(400, 200),
            nn.ReLU(inplace=True),
            nn.Linear(200, 10)
        )

    def forward(self, x):

        # TODO: define your model here
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1,400)
        x = self.fc_net(x)

        return x
```

1.2 Print command `print(net)`

```
BaseNet(  
  (conv_net): Sequential(  
    (0): Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))  
    (1): ReLU(inplace=True)  
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (3): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))  
    (4): ReLU(inplace=True)  
    (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (fc_net): Sequential(  
    (0): Linear(in_features=400, out_features=200, bias=True)  
    (1): ReLU(inplace=True)  
    (2): Linear(in_features=200, out_features=10, bias=True)  
  )  
)
```

1.3 Final accuracy for train the model

Final accuracy on the validation set

```
Accuracy of the final network on the val images: 59.9 %  
Accuracy of airplane : 62.9 %  
Accuracy of automobile : 73.6 %  
Accuracy of bird : 41.4 %  
Accuracy of cat : 40.7 %  
Accuracy of deer : 58.2 %  
Accuracy of dog : 49.8 %  
Accuracy of frog : 69.7 %  
Accuracy of horse : 71.1 %  
Accuracy of ship : 65.5 %  
Accuracy of truck : 66.4 %
```

```
[1] loss: 1.867  
Accuracy of the network on the val images: 44 %  
[2] loss: 1.461  
Accuracy of the network on the val images: 50 %  
[3] loss: 1.327  
Accuracy of the network on the val images: 55 %  
[4] loss: 1.232  
Accuracy of the network on the val images: 56 %  
[5] loss: 1.146  
Accuracy of the network on the val images: 56 %  
[6] loss: 1.082
```

```
Accuracy of the network on the val images: 60 %  
[7] loss: 1.024  
Accuracy of the network on the val images: 61 %  
[8] loss: 0.968  
Accuracy of the network on the val images: 60 %  
[9] loss: 0.925  
Accuracy of the network on the val images: 60 %  
[10] loss: 0.875  
Accuracy of the network on the val images: 60 %  
[11] loss: 0.834  
Accuracy of the network on the val images: 61 %  
[12] loss: 0.798  
Accuracy of the network on the val images: 61 %  
[13] loss: 0.764  
Accuracy of the network on the val images: 60 %  
[14] loss: 0.729  
Accuracy of the network on the val images: 59 %  
[15] loss: 0.700  
Accuracy of the network on the val images: 59 %  
  
Finished Training
```

2. Improve BaseNet

My best model Code:

2.1 Final accuracy on validation set

```
Accuracy of the final network on the val images: 86.9 %  
Accuracy of airplane : 89.5 %  
Accuracy of automobile : 94.3 %  
Accuracy of bird : 77.2 %  
Accuracy of cat : 73.3 %  
Accuracy of deer : 88.5 %  
Accuracy of dog : 76.8 %  
Accuracy of frog : 93.3 %  
Accuracy of horse : 92.4 %  
Accuracy of ship : 90.6 %  
Accuracy of truck : 92.7 %
```


2.2 Table defining your final architecture

| Layer No. | Layer Type No. | Layer Type | Kernel Size | Input Dim | Output Dim | Input Channels | Output Channels |
|-----------|----------------|-------------|-------------|-----------|------------|----------------|-----------------|
| 1 | Conv2d 1 | Conv2d | 3 | 32 | 32 | 3 | 64 |
| 2 | BatchNorm2d | BatchNorm2d | - | 32 | 32 | 64 | 64 |
| 3 | Relu | Relu | - | 32 | 32 | 64 | 64 |
| 4 | Conv2d 2 | Conv2d | 3 | 32 | 32 | 64 | 64 |
| 5 | BatchNorm2d | BatchNorm2d | - | 32 | 32 | 64 | 64 |
| 6 | Relu | Relu | - | 32 | 32 | 64 | 64 |
| 7 | Conv2d 3 | Conv2d | 3 | 32 | 32 | 64 | 64 |
| 8 | BatchNorm2d | BatchNorm2d | - | 32 | 32 | 64 | 64 |
| 9 | Relu | Relu | - | 32 | 32 | 64 | 64 |
| 10 | Conv2d 4 | Conv2d | 3 | 32 | 32 | 64 | 64 |
| 11 | BatchNorm2d | BatchNorm2d | - | 32 | 32 | 64 | 64 |
| 12 | Relu | Relu | - | 32 | 32 | 64 | 64 |
| 13 | Conv2d 5 | Conv2d | 3 | 32 | 32 | 64 | 64 |
| 14 | BatchNorm2d | BatchNorm2d | - | 32 | 32 | 64 | 64 |
| 15 | Relu | Relu | - | 32 | 32 | 64 | 64 |
| 16 | Conv2d 6 | Conv2d | 3 | 32 | 32 | 64 | 128 |
| 17 | BatchNorm2d | BatchNorm2d | - | 32 | 32 | 128 | 128 |
| 18 | Relu | Relu | - | 32 | 32 | 128 | 128 |
| 19 | Conv2d 7 | Conv2d | 3 | 32 | 32 | 128 | 128 |
| 20 | BatchNorm2d | BatchNorm2d | - | 32 | 32 | 128 | 128 |
| 21 | Relu | Relu | - | 32 | 32 | 128 | 128 |
| 22 | Maxpool 1 | Maxpool | 2 | 32 | 16 | 128 | 128 |
| 23 | Conv2d 8 | Conv2d | 3 | 16 | 16 | 128 | 128 |
| 24 | BatchNorm2d | BatchNorm2d | - | 16 | 16 | 128 | 128 |
| 25 | Relu | Relu | - | 16 | 16 | 128 | 128 |
| 26 | Conv2d 9 | Conv2d | 3 | 16 | 16 | 128 | 128 |
| 27 | BatchNorm2d | BatchNorm2d | - | 16 | 16 | 128 | 128 |
| 28 | Relu | Relu | - | 16 | 16 | 128 | 128 |
| 29 | Conv2d 10 | Conv2d | 3 | 16 | 16 | 128 | 128 |
| 30 | BatchNorm2d | BatchNorm2d | - | 16 | 16 | 128 | 128 |
| 31 | Relu | Relu | - | 16 | 16 | 128 | 128 |
| 32 | Conv2d 11 | Conv2d | 3 | 16 | 16 | 128 | 128 |
| 33 | BatchNorm2d | BatchNorm2d | - | 16 | 16 | 128 | 128 |
| 34 | Relu | Relu | - | 16 | 16 | 128 | 128 |
| 35 | Conv2d 12 | Conv2d | 3 | 16 | 16 | 128 | 256 |
| 36 | BatchNorm2d | BatchNorm2d | - | 16 | 16 | 256 | 256 |
| 37 | Relu | Relu | - | 16 | 16 | 256 | 256 |

| | | | | | | | |
|----|-------------|-------------|---|----|----|------|------|
| 38 | Conv2d 13 | Conv2d | 3 | 16 | 16 | 256 | 256 |
| 39 | BatchNorm2d | BatchNorm2d | - | 16 | 16 | 256 | 256 |
| 40 | Relu | Relu | - | 16 | 16 | 256 | 256 |
| 41 | Maxpool 2 | Maxpool | 2 | 16 | 8 | 256 | 256 |
| 42 | Conv2d 14 | Conv2d | 3 | 8 | 8 | 256 | 256 |
| 43 | BatchNorm2d | BatchNorm2d | - | 8 | 8 | 256 | 256 |
| 44 | Relu | Relu | - | 8 | 8 | 256 | 256 |
| 45 | Conv2d 15 | Conv2d | 3 | 8 | 8 | 256 | 256 |
| 46 | BatchNorm2d | BatchNorm2d | - | 8 | 8 | 256 | 256 |
| 47 | Relu | Relu | - | 8 | 8 | 256 | 256 |
| 48 | Conv2d 16 | Conv2d | 3 | 8 | 8 | 256 | 256 |
| 49 | BatchNorm2d | BatchNorm2d | - | 8 | 8 | 256 | 256 |
| 50 | Relu | Relu | - | 8 | 8 | 256 | 256 |
| 51 | Conv2d 17 | Conv2d | 3 | 8 | 8 | 256 | 256 |
| 52 | BatchNorm2d | BatchNorm2d | - | 8 | 8 | 256 | 256 |
| 53 | Relu | Relu | - | 8 | 8 | 256 | 256 |
| 54 | Conv2d 18 | Conv2d | 3 | 8 | 8 | 256 | 512 |
| 55 | BatchNorm2d | BatchNorm2d | - | 8 | 8 | 512 | 512 |
| 56 | Relu | Relu | - | 8 | 8 | 512 | 512 |
| 57 | Conv2d 19 | Conv2d | 3 | 8 | 8 | 512 | 512 |
| 58 | BatchNorm2d | BatchNorm2d | - | 8 | 8 | 512 | 512 |
| 59 | Relu | Relu | - | 8 | 8 | 512 | 512 |
| 60 | Maxpool 3 | Maxpool | 2 | 8 | 4 | 512 | 512 |
| 61 | Conv2d 20 | Conv2d | 3 | 4 | 4 | 512 | 512 |
| 62 | BatchNorm2d | BatchNorm2d | - | 4 | 4 | 512 | 512 |
| 63 | Relu | Relu | - | 4 | 4 | 512 | 512 |
| 64 | Conv2d 21 | Conv2d | 3 | 4 | 4 | 512 | 512 |
| 65 | BatchNorm2d | BatchNorm2d | - | 4 | 4 | 512 | 512 |
| 66 | Relu | Relu | - | 4 | 4 | 512 | 512 |
| 67 | Conv2d 22 | Conv2d | 3 | 4 | 4 | 512 | 512 |
| 68 | BatchNorm2d | BatchNorm2d | - | 4 | 4 | 512 | 512 |
| 69 | Relu | Relu | - | 4 | 4 | 512 | 512 |
| 70 | Conv2d 23 | Conv2d | 3 | 4 | 4 | 512 | 512 |
| 71 | BatchNorm2d | BatchNorm2d | - | 4 | 4 | 512 | 512 |
| 72 | Relu | Relu | - | 4 | 4 | 512 | 512 |
| 73 | Conv2d 24 | Conv2d | 3 | 4 | 4 | 512 | 1024 |
| 74 | BatchNorm2d | BatchNorm2d | - | 4 | 4 | 1024 | 1024 |
| 75 | Relu | Relu | - | 4 | 4 | 1024 | 1024 |
| 76 | Conv2d 25 | Conv2d | 3 | 4 | 4 | 1024 | 1024 |
| 77 | BatchNorm2d | BatchNorm2d | - | 4 | 4 | 1024 | 1024 |

| | | | | | | | |
|----|-----------|---------|---|---|---|------|------|
| 78 | Relu | Relu | - | 4 | 4 | 1024 | 1024 |
| 79 | Maxpool 4 | Maxpool | 2 | 4 | 2 | 1024 | 1024 |
| 80 | Linear | Linear | - | 1 | 1 | 1024 | 512 |
| 81 | Relu | Relu | - | 1 | 1 | 512 | 512 |
| 82 | Linear | Linear | - | 1 | 1 | 512 | 10 |
| | | | | | | | |

2.3 Training loss plot and test accuracy plot for final model



2.4 An ablation table, listing all factors that you tried to make improvement to your final model as well as the corresponding validation accuracy.

| Improvement | Detail | performance |

| Improvement | Detail | performance |
|-----------------------|--|-------------------|
| Data normalization | : Normalizing input data makes training easier and more robust. Normalize the data to made it zero mean and fixed standard deviation. So I use the same transforms.Normalize() for both train and test dataset. | 78.89% |
| Data augmentation | Augment the training data using random crops, horizontal flips, etc. I use transforms.RandomHorizontalFlip() for my train dataset. train_transform = transforms.Compose([transforms.RandomHorizontalFlip(), transforms.ToTensor(), transforms.Normalize(mean = (0.5,), std = (0.5,)),]) test_transform = transforms.Compose([transforms.ToTensor() , transforms.Normalize(mean = (0.5,), std = (0.5,)),]) | 78.89% |
| Deeper network. | I added a lot of convolutional layers for training to build a deeper network. And end up having an model with around 80-90 layers. | 70.23%, 80.93% |
| Normalization layers. | For each of them, I also added the normalization layers after conv layers (nn.BatchNorm2d). | 78.87% |
| Change optimizer | I tried optimizer like SGD and Adam. | 83.89% |
| Change learning rate | I tried the learning rate of 0.005 and 0.01 | 83.71% |
| Increase epochs | I found the accuracty is still increasing, so I increase the epochs to 70, and the accuracy is almost stable at that time | 70.23% |

3. Secret test set

Autograder Output

```
Q1 evaluation results:
Accuracy: 77.5 %
Accuracy of airplane : 77.0 %
Accuracy of automobile : 72.5 %
Accuracy of bird : 71.5 %
Accuracy of cat : 65.0 %
Accuracy of deer : 80.5 %
Accuracy of dog : 67.5 %
Accuracy of frog : 83.5 %
Accuracy of horse : 85.5 %
Accuracy of ship : 81.5 %
Accuracy of truck : 91.0 %
```

```
100%|██████████| 200/200 [00:15<00:00, 12.69it/s]
Q2 evaluation results:
Test metrics:
mean error 24.8
median error 16.2
accuracy at 11.25deg 38.3
accuracy at 22.5deg 61.5
accuracy at 30deg 71.0
```

```
=====
Results found for Q1, good work!
=====
```

```
=====
Results found for Q2, good work!
=====
```

Problem2 Surface normal

2.Build on top of ImageNet pre-trained Model

`print(net)`

```
MyModel(
  (layers): ModuleList(
    (0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
    (4): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
```



```

(4): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
  (1): BasicBlock(
    (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
(5): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)

```

```

        (relu): ReLU(inplace=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(6): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
(7): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,

```

```


        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (downsample): Sequential(
          (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
          (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    (1): BasicBlock(
      (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (8): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1))
  (9): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (10): ReLU(inplace=True)
  (11): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1))
  (12): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (13): ReLU(inplace=True)
  (14): Conv2d(128, 64, kernel_size=(1, 1), stride=(1, 1))
  (15): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (16): ReLU(inplace=True)
  (17): Conv2d(64, 3, kernel_size=(1, 1), stride=(1, 1))
  (18): Upsample(size=(512, 512), mode=bilinear)
)
)

```

Final performance on validation set (all 5 metrics)

| mean angular error (lower the better) | median angular error (lower the better) | accuracies at 11.25 degree (higher the better) | accuracies at 22.5 degree (higher the better) | accuracies at 30 degree (higher the better) |
|---------------------------------------|---|--|---|---|
| 34.1 | 26.7 | 23.6 | 44.4 | 53.9 |

Validation loss (L1): 0.2499153511837507
Validation metrics: Mean 34.1, Median 26.7, 11.25deg 23.6, 22.5deg 44.4, 30deg 53.9



3. Increase your model output resolution


Your best model. Include final performance on validation set (5 metrics).

Final performance on validation set (all 5 metrics)

| mean angular error (lower the better) | median angular error (lower the better) | accuracies at 11.25 degree (higher the better) | accuracies at 22.5 degree (higher the better) | accuracies at 30 degree (higher the better) |
|---------------------------------------|---|--|---|---|
| 26.9 | 18.6 | 36.2 | 56.0 | 65.5 |

Validation loss (L1): 0.1972911968254126

Validation metrics: Mean 26.9, Median 18.6, 11.25deg 36.2, 22.5deg 56.0, 30deg 65.5



An ablation table, listing all factors that you tried to make improvement to your final model as well as the validation performance.

I have tried several different models.

| Models and methods | details | validation performance |
|--|--|---|
| Improve on Resnet18 | I added several Relu() and Batch() layers, | Validation metrics: Mean 32.7, Median 26.8, 11.25deg 23.2, 22.5deg 43.6, 30deg 54.4 |
| Unet | I use the U-net | Validation metrics: Mean 34.8, Median 25.4, 11.25deg 26.1, 22.5deg 46.2, 30deg 55.3 |
| DeepLabv3+ + Adam optimizer | I use this optimizer torch.optim.Adam(model.parameters(),lr=0.005) | Validation metrics: Mean 30.2, Median 22.8, 11.25deg 27.7, 22.5deg 49.6, 30deg 60.0 |
| DeepLabv3+ + SGD optimizer | I use this optimizer optim.SGD(model.parameters(), lr=0.005, momentum=0.9, weight_decay=0.001) | Validation metrics: Mean 27.4, Median 19.4, 11.25deg 32.8, 22.5deg 55.0, 30deg 65.0 |
| Increase epochs | For the Unet, I have tried to use epochs as large as 100. For the Deeplabv3+, I use around 30. | Validation metrics: Mean 34.8, Median 25.4, 11.25deg 26.1, 22.5deg 46.2, 30deg 55.3 |
| Change learning rate | I tried 0.005 and 0.01 | Validation metrics: Mean 30.2, Median 22.8, 11.25deg 27.7, 22.5deg 49.6, 30deg 60.0 |
| Change weight decay | I tried weight decay of 0, 0.001 to 0.0. Finally, I use 0.001 | Validation metrics: Mean 34.8, Median 25.4, 11.25deg 26.1, 22.5deg 46.2, 30deg 55.3 |
| Tried different activation layers | Tried different activation layers: Relu(), sigmoid | Validation metrics: Mean 34.8, Median 25.4, 11.25deg 26.1, 22.5deg 46.2, 30deg 55.3 |
| Tried normalization layers | Tried normalization layers: such as batchNorm2d().after conv() layers | Validation metrics: Mean 34.8, Median 25.4, 11.25deg 26.1, 22.5deg 46.2, 30deg 55.3 |
| Add more layers and build a deeper network | I start with the simple one of res18 and I got a very complex network at the end. | Validation metrics: Mean 30.2, Median 22.8, 11.25deg 27.7, 22.5deg 49.6, 30deg 60.0 |
| Tried different ways of upsample. | Like, upsample once of scale 32 or upsample scale 2 for several times. | Validation metrics: Mean 30.2, Median 22.8, 11.25deg 27.7, 22.5deg 49.6, 30deg 60.0 |

Also, for the main model, I also have the model store the current best performance, in case we will overfit,

More Description for my final model.

I implement the a DeepLabv3+ decoder.

My optimizer is SGD with learning rate of 0.005

```
optimizer= optim.SGD(model.parameters(), lr=0.005, momentum=0.9, weight_decay=0.001)
```

And my loss function is

```
loss = (F.l1_loss(prediction, target, reduction = "none") * mask).nansum()/512/512
```

My model now looks like this:

```
Using cache found in /root/.cache/torch/hub/pytorch_vision_v0.9.0
MyModel(
  (deeplab): DeepLabV3(
    (backbone): IntermediateLayerGetter(
      (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
ceil_mode=False)
      (layer1): Sequential(
        (0): Bottleneck(
          (conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
          (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (relu): ReLU(inplace=True)
          (downsample): Sequential(
            (0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
            (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          )
        )
      (1): Bottleneck(
        (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
        (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
```

```

    )
    (2): Bottleneck(
      (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
  )
  (layer2): Sequential(
    (0): Bottleneck(
      (conv1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (downsample): Sequential(
        (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (1): Bottleneck(
      (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)

```

```

    )
    (2): Bottleneck(
      (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (3): Bottleneck(
      (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
  )
  (layer3): Sequential(
    (0): Bottleneck(
      (conv1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (downsample): Sequential(
        (0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
  )

```



```

    )
    (1): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (2): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (3): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (4): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),

```

```

dilation=(2, 2), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
)
(5): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
)
(6): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
)
(7): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
)

```

```

    )
    (8): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (9): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (10): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (11): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),

```

```

dilation=(2, 2), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
)
(12): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
)
(13): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
)
(14): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
)

```

```

    )
    (15): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (16): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (17): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (18): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),

```

```

dilation=(2, 2), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
)
(19): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
)
(20): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
)
(21): Bottleneck(
    (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
)

```

```

    )
    (22): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
  )
  (layer4): Sequential(
    (0): Bottleneck(
      (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2), bias=False)
      (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (downsample): Sequential(
        (0): Conv2d(1024, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      )
    )
    (1): Bottleneck(
      (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(4, 4),
dilation=(4, 4), bias=False)
      (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
  )

```

```

    )
    (2): Bottleneck(
      (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(4, 4),
dilation=(4, 4), bias=False)
      (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
  )
)
(classifier): DeepLabHead(
  (0): ASPP(
    (convs): ModuleList(
      (0): Sequential(
        (0): Conv2d(2048, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU()
      )
      (1): ASPPConv(
        (0): Conv2d(2048, 256, kernel_size=(3, 3), stride=(1, 1), padding=(12, 12),
dilation=(12, 12), bias=False)
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU()
      )
      (2): ASPPConv(
        (0): Conv2d(2048, 256, kernel_size=(3, 3), stride=(1, 1), padding=(24, 24),
dilation=(24, 24), bias=False)
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU()
      )
      (3): ASPPConv(
        (0): Conv2d(2048, 256, kernel_size=(3, 3), stride=(1, 1), padding=(36, 36),
dilation=(36, 36), bias=False)
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU()
      )
    )
  )
)

```



```


        (4): ASPPPooling(
          (0): AdaptiveAvgPool2d(output_size=1)
          (1): Conv2d(2048, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
          (2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (3): ReLU()
        )
      )
    (project): Sequential(
      (0): Conv2d(1280, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): ReLU()
      (3): Dropout(p=0.5, inplace=False)
    )
  )
  (1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (3): ReLU()
  (4): Conv2d(256, 3, kernel_size=(1, 1), stride=(1, 1))
)
(aux_classifier): FCNHead(
  (0): Conv2d(1024, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
  (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (2): ReLU()
  (3): Dropout(p=0.1, inplace=False)
  (4): Conv2d(256, 21, kernel_size=(1, 1), stride=(1, 1))
)
)
)
)

```


4. Visualize your prediction

Visual comparisons of the output from part 2

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



Visual comparisons of the output from part 3



Visual comparisons of the output from part 2 and 3 on 2 images from the validation dataset, discuss your observations.

- We can see that the output from part3 is more smooth than the part2
- Also, we can see that the differences between different part of the pictures is most clear, have a better contour for part 3.
- We can see more details in the pictures generated by part 3.
- The output of part3 is very close to the original pictures.

5. Secret test set

Autograder Output

Q1 evaluation results:
Accuracy: 77.5 %
Accuracy of airplane : 77.0 %
Accuracy of automobile : 72.5 %
Accuracy of bird : 71.5 %
Accuracy of cat : 65.0 %
Accuracy of deer : 80.5 %
Accuracy of dog : 67.5 %
Accuracy of frog : 83.5 %
Accuracy of horse : 85.5 %
Accuracy of ship : 81.5 %
Accuracy of truck : 91.0 %

100% ██████████ 200/200 [00:15<00:00, 12.69it/s]
Q2 evaluation results:
Test metrics:
mean error 24.8
median error 16.2
accuracy at 11.25deg 38.3
accuracy at 22.5deg 61.5
accuracy at 30deg 71.0

Results found for Q1, good work!

Results found for Q2, good work!

AUTOGRADER SCORE
10.0 / 10.0

| mean angular error (lower the better) | median angular error (lower the better) | accuracies at 11.25 degree (higher the better) | accuracies at 22.5 degree (higher the better) | accuracies at 30 degree (higher the better) |
|---------------------------------------|---|--|---|---|
| 24.8 | 16.2 | 38.3 | 61.5 | 71.0 |

Rank as top 1 so far.

| Rank | Submission Name | Q1: Overall Accuracy | Q2: Mean Angular Error | Q2: Median Angular Error | Q2: Accuracy at 11.25 Deg | Q2: Accuracy at 22.5 Deg | Q2: Accuracy at 30 Deg |
|------|-----------------|----------------------|------------------------|--------------------------|---------------------------|--------------------------|------------------------|
| 93 | xxx | 77.55 | 24.8 | 16.18 | 38.28 | 61.54 | 71.03 |
| 92 | Romo | 0 | 27.29 | 17.99 | 36.05 | 56.96 | 66.16 |
| 91 | Mishra | 70.6 | 27.55 | 16.96 | 36.75 | 60.74 | 67.57 |